

# Exercise 1

*Ken Benoit*

*26 September 2014*

## Using quanteda in R

This take-home exercise is designed to get you working with quanteda (<http://github.com/kbenoit/quanteda>). The focus will be on exploring the package and getting some texts into the **corpus** object format. quanteda (<http://github.com/kbenoit/quanteda>) package has several functions for creating a corpus of texts which we will use in this exercise.

### 1. Getting Started.

You can use R or Rstudio for these exercises. You will first need to install the package, using:

```
# needs the devtools package for this to work
if (!require(devtools)) install.packages("devtools", dependencies=TRUE)
# be sure to install the latest version from GitHub, using dev branch:
devtools::install_github("kbenoit/quanteda", username="kbenoit", dependencies=TRUE, ref="dev")
# and quantedaData
devtools::install_github("kbenoit/quantedaData", username="kbenoit")
```

### 2. Exploring **quanteda** functions.

You can try running `demo(quanteda)`, and also use the `example()` function for any function in the package, to run the examples and see how the function works. Of course you should also browse the documentation, especially `?corpus` to see the structure and operations of how to construct a corpus.

### 3. Making a corpus and corpus structure

#### 1. From a vector of texts already in memory.

The simplest way to create a corpus is to use a vector of texts already present in R's global environment. Some text and corpus objects are built into the package, for example `inaugTexts` is the UTF-8 encoded set of 57 presidential inaugural addresses. Try using `corpus()` on this set of texts to create a corpus.

Once you have constructed this corpus, use the `summary()` method to see a brief description of the corpus. The names of the character vector `inaugTexts` should have become the document names.

#### 2. From a directory of text files.

The `corpus()` function can take as its main argument the name of a directory, if you wrap the path to the directory within a `directory()` call. (See `?directory` for an example.) If you call `directory()` with no arguments, then it should allow you to choose the directory interactively (you will need to have installed the `tc1tk2` package first though.)

Here you are encouraged to select any directory of plain text files of your own.  
How did it work? Try using `docvars()` to assign a set of document-level variables.

Note that if you document level metadata in your filenames, then this can be automatically parsed by `corpus.directory()` into `docvars`.

```
require(quantda)
```

```
## Loading required package: quantda
```

```
mydir <- directory("~/Dropbox/QUANTESS/corpora/ukManRenamed")
mycorpus <- corpus(mydir)
summary(mycorpus, 5)
```

```
## Corpus consisting of 101 documents.
##
##           Text Types Tokens Sentences docvar1 docvar2 docvar3
## UK_natl_1945_en_Con 1578   6095      275      UK   natl   1945
## UK_natl_1945_en_Lab 1258   4975      241      UK   natl   1945
## UK_natl_1945_en_Lib 1060   3377      158      UK   natl   1945
## UK_natl_1950_en_Con 1800   7413      381      UK   natl   1950
## UK_natl_1950_en_Lab 1342   4879      275      UK   natl   1950
## docvar4 docvar5
##      en      Con
##      en      Lab
##      en      Lib
##      en      Con
##      en      Lab
##
## Source: /Users/kbenoit/Dropbox/Classes/QTA NYU/Exercises/Exercise 1/* on x86_64 by kbenoi
t.
## Created: Fri Sep 26 12:39:15 2014.
## Notes: .
```

3. From a zipped file. For this, you can try downloading the zipped file of Irish budget speeches available from <http://www.kenbenoit.net/courses/nyu2014qta/iebudget2010.zip> (<http://www.kenbenoit.net/courses/nyu2014qta/iebudget2010.zip>). Then execute this code:

```
myzipfiles <- zipfiles() # allows you to locate the file interactively
mycorpus <- corpus(myzipfiles,
                   docvarnames=c("year", "debate", "seq", "fname", "lname", "party"))
summary(mycorpus, 5)
```

4. If you are familiar with the **tm** text package for R, it is also now possible to import one of its `VCorpus` objects directly, using the `corpus.VCorpus()` method.  
This is a feature I just added, so I am eager for you to test it.

#### 4. Explore some phrases in the text.

You can do this using the `kwic` (for “key-words-in-context”) to explore a specific word or phrase.

```
kwic(inaugCorpus, "terror", 3)
```

```
##                                preword      word
##      [1797-Adams, 1183]           or violence, by  terror,
## [1933-Roosevelt, 100] nameless, unreasoning, unjustified  terror
## [1941-Roosevelt, 252]           by a fatalistic  terror,
##      [1961-Kennedy, 763]         uncertain balance of  terror
##      [1961-Kennedy, 872]         instead of its  terrors.
##      [1981-Reagan, 691]         Americans from the  terror
##      [1981-Reagan, 1891]        those who practice terrorism
##      [1997-Clinton, 929]        the fanaticism of  terror.
## [1997-Clinton, 1462]        strong defense against  terror
##      [2009-Obama, 1433]        aims by inducing  terror
##                                postword
##      [1797-Adams, 1183] intrigue, or venality,
## [1933-Roosevelt, 100] which paralyzes needed
## [1941-Roosevelt, 252] we proved that
##      [1961-Kennedy, 763] that stays the
##      [1961-Kennedy, 872] Together let us
##      [1981-Reagan, 691] of runaway living
##      [1981-Reagan, 1891] and prey upon
##      [1997-Clinton, 929] And they torment
## [1997-Clinton, 1462] and destruction. Our
##      [2009-Obama, 1433] and slaughtering innocents,
```

Try substituting your own search terms, or working with your own corpus.

#### 5. Create a document-feature matrix, using `dfm`. First, read the documentation using `?dfm` to see the available options.

```
mydfm <- dfm(inaugCorpus, stopwords=TRUE)
```

```
## Creating dfm from a corpus: ... removing stopwords ... done.
```

```
dim(mydfm)
```

```
## [1]    57 9087
```

```
topfeatures(mydfm, 20)
```

##	will	people	government	us	can	upon
##	871	564	561	476	470	371
##	must	may	great	states	shall	world
##	363	338	334	331	314	305
##	country	every	nation	peace	one	new
##	294	291	287	253	244	241
##	power	public				
##	232	223				

Experiment with different `dfm` options, such as `stem=TRUE`. The function `trimdfm()` allows you to reduce the size of the dfm following its construction.

Grouping on a variable is an excellent feature of `dfm()`, in fact one of my favorites. For instance, if you want to aggregate all speeches by presidential name, you can execute

```
mydfm <- dfm(inaugCorpus, groups="President")
```

```
## Creating dfm from a corpus: ... aggregating by group: President... complete ... done.
```

```
dim(mydfm)
```

```
## [1] 34 9210
```

```
docnames(mydfm)
```

```
## [1] "Adams"      "Buchanan"   "Bush"       "Carter"     "Cleveland"
## [6] "Clinton"   "Coolidge"   "Eisenhower" "Garfield"   "Grant"
## [11] "Harding"    "Harrison"   "Hayes"      "Hoover"     "Jackson"
## [16] "Jefferson"  "Johnson"   "Kennedy"    "Lincoln"    "Madison"
## [21] "McKinley"   "Monroe"     "Nixon"      "Obama"      "Pierce"
## [26] "Polk"       "Reagan"     "Roosevelt"  "Taft"       "Taylor"
## [31] "Truman"     "VanBuren"   "Washington" "Wilson"
```

Note that this groups Theodore and Franklin D. Roosevelt together – to separate them we would have needed to add a `firstname` variable using `docvars()` and grouped on that as well.

## 6. Explore the ability to subset a corpus.


There is a `subset()` method defined for a corpus, which works just like R's normal `subset()` command. This provides an easy method to send specific documents to downstream functions, like `dfm()`, which will be useful workaround until I implement a full set of subsetting and indexing features for the `dfm` class object.

For instance if you want a wordcloud of just Obama's two inaugural addresses, you would need to subset the corpus first:

```
obamadfm <- dfm(subset(inaugCorpus, President=="Obama"), stopwords=TRUE)
```

```
## Creating dfm from a corpus: ... removing stopwords ... done.
```

```
plot(obamadfm)
```



## Bug reports

**quanteda** is a work in progress. Please send me suggestions, bug reports, etc. so that I can improve it. You can email these directly to [kbenoit@lse.ac.uk](mailto:kbenoit@lse.ac.uk) (<mailto:kbenoit@lse.ac.uk>).

In addition, if you are having trouble with importing your texts into a corpus, I welcome you to send me a set of your texts and I will write the code for you (and possibly add to the functionality of **quanteda** to make this possible). I am particularly interested in: \* alternative formats - csv, pdf, Word, XML, etc. \* non-English languages - encoding is one of the issues we plan to tackle \* interesting/non-standard document units \* large volumes of text, to see **quanteda**'s functionality scales.